

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [2]: loans = pd.read_csv('loan_data.csv')
```

```
In [3]: loans.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9578 entries, 0 to 9577
Data columns (total 14 columns):
credit.policy      9578 non-null int64
purpose           9578 non-null object
int.rate          9578 non-null float64
installment       9578 non-null float64
log.annual.inc    9578 non-null float64
dti               9578 non-null float64
fico              9578 non-null int64
days.with.cr.line 9578 non-null float64
revol.bal         9578 non-null int64
revol.util        9578 non-null float64
inq.last.6mths    9578 non-null int64
delinq.2yrs       9578 non-null int64
pub.rec           9578 non-null int64
not.fully.paid    9578 non-null int64
dtypes: float64(6), int64(7), object(1)
memory usage: 1.0+ MB
```

```
In [4]: loans.describe()
```

Out[4]:

credit.policy	int.rate	installment	log.annual.inc	dti	fico	days.witl
---------------	----------	-------------	----------------	-----	------	-----------

	credit.policy	int.rate	installment	log.annual.inc	dti	fico	days.with
count	9578.000000	9578.000000	9578.000000	9578.000000	9578.000000	9578.000000	9578
mean	0.804970	0.122640	319.089413	10.932117	12.606679	710.846314	4560
std	0.396245	0.026847	207.071301	0.614813	6.883970	37.970537	2496
min	0.000000	0.060000	15.670000	7.547502	0.000000	612.000000	178
25%	1.000000	0.103900	163.770000	10.558414	7.212500	682.000000	2820
50%	1.000000	0.122100	268.950000	10.928884	12.665000	707.000000	4139
75%	1.000000	0.140700	432.762500	11.291293	17.950000	737.000000	5730
max	1.000000	0.216400	940.140000	14.528354	29.960000	827.000000	17639


In [5]: `loans.head()`

Out[5]:

	credit.policy	purpose	int.rate	installment	log.annual.inc	dti	fico	days.with.cr.line
0	1	debt_consolidation	0.1189	829.10	11.350407	19.48	737	5639.958333
1	1	credit_card	0.1071	228.22	11.082143	14.29	707	2760.000000
2	1	debt_consolidation	0.1357	366.86	10.373491	11.63	682	4710.000000
3	1	debt_consolidation	0.1008	162.34	11.350407	8.10	712	2699.958333
4	1	credit_card	0.1426	102.92	11.299732	14.97	667	4066.000000


```
In [6]: plt.figure(figsize=(10,6))
        loans[loans['credit.policy']==1]['fico'].hist(alpha=0.5,color='blue',
                                                    bins=30,label='Credit.Policy=1')
        loans[loans['credit.policy']==0]['fico'].hist(alpha=0.5,color='red',
                                                    bins=30,label='Credit.Policy=0')
        plt.legend()
        plt.xlabel('FICO')
```

Out[6]: Text(0.5,0,'FICO')




```
In [7]: plt.figure(figsize=(10,6))
        loans[loans['not.fully.paid']==1]['fico'].hist(alpha=0.5,color='blue',
                                                    bins=30,label='not.fully.
        paid=1')
        loans[loans['not.fully.paid']==0]['fico'].hist(alpha=0.5,color='red',
                                                    bins=30,label='not.fully.
        paid=0')
        plt.legend()
        plt.xlabel('FICO')
```

Out[7]: Text(0.5,0,'FICO')



```
In [8]: plt.figure(figsize=(11,7))  
sns.countplot(x='purpose',hue='not.fully.paid',data=loans,palette='Set  
1')
```

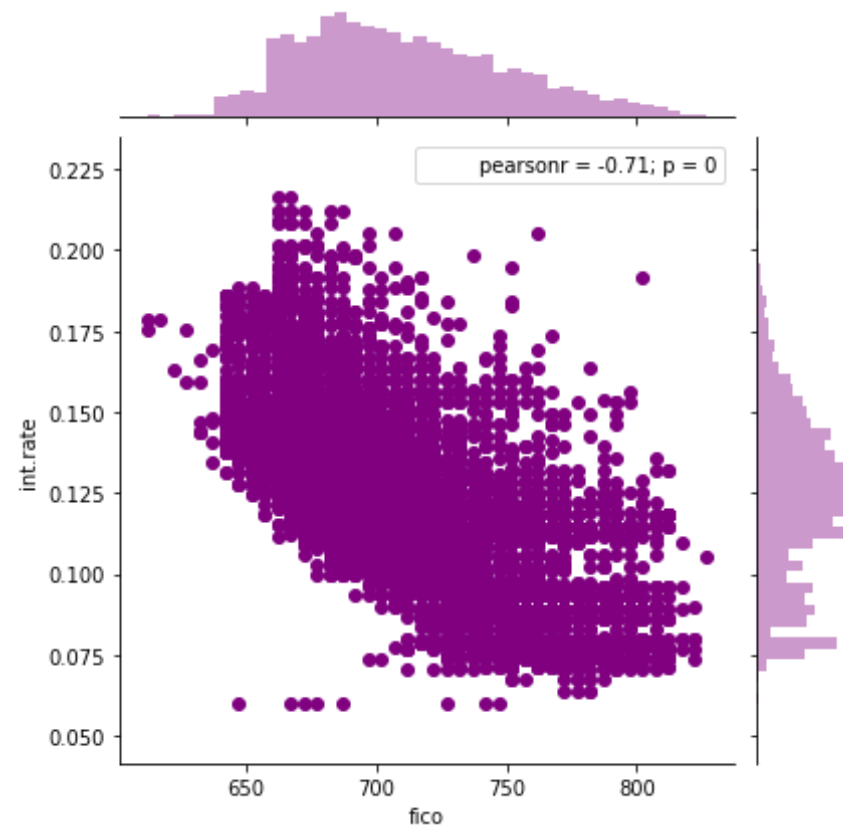
```
Out[8]: <matplotlib.axes._subplots.AxesSubplot at 0x106921b38>
```



```
In [9]: sns.jointplot(x='fico',y='int.rate',data=loans,color='purple')
```


```
/anaconda3/lib/python3.6/site-packages/matplotlib/axes/_axes.py:6462: U
serWarning: The 'normed' kwarg is deprecated, and has been replaced by
the 'density' kwarg.
  warnings.warn("The 'normed' kwarg is deprecated, and has been "
/anaconda3/lib/python3.6/site-packages/matplotlib/axes/_axes.py:6462: U
serWarning: The 'normed' kwarg is deprecated, and has been replaced by
the 'density' kwarg.
  warnings.warn("The 'normed' kwarg is deprecated, and has been "
```

```
Out[9]: <seaborn.axisgrid.JointGrid at 0x1a0cc92550>
```



```
In [10]: plt.figure(figsize=(11,7))
sns.lmplot(y='int.rate',x='fico',data=loans,hue='credit.policy',
          col='not.fully.paid',palette='Set1')
```

```
Out[10]: <seaborn.axisgrid.FacetGrid at 0x1a18d922e8>
<Figure size 792x504 with 0 Axes>
```



In [11]: `loans.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9578 entries, 0 to 9577
Data columns (total 14 columns):
credit.policy      9578 non-null int64
purpose           9578 non-null object
int.rate          9578 non-null float64
installment       9578 non-null float64
log.annual.inc    9578 non-null float64
dti               9578 non-null float64
fico              9578 non-null int64
days.with.cr.line 9578 non-null float64
revol.bal         9578 non-null int64
revol.util        9578 non-null float64
inq.last.6mths    9578 non-null int64
delinq.2yrs       9578 non-null int64
pub.rec           9578 non-null int64
not.fully.paid    9578 non-null int64
dtypes: float64(6), int64(7), object(1)
memory usage: 1.0+ MB
```

```
In [12]: cat_feats = ['purpose']
```

```
In [13]: final_data = pd.get_dummies(loans, columns=cat_feats, drop_first=True)
```

```
In [14]: final_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 9578 entries, 0 to 9577  
Data columns (total 19 columns):  
credit.policy          9578 non-null int64  
int.rate              9578 non-null float64  
installment           9578 non-null float64  
log.annual.inc        9578 non-null float64  
dti                   9578 non-null float64  
fico                  9578 non-null int64  
days.with.cr.line    9578 non-null float64  
revol.bal             9578 non-null int64  
revol.util            9578 non-null float64  
inq.last.6mths        9578 non-null int64  
delinq.2yrs           9578 non-null int64  
pub.rec               9578 non-null int64  
not.fully.paid         9578 non-null int64  
purpose_credit_card   9578 non-null uint8  
purpose_debt_consolidation 9578 non-null uint8  
purpose_educational   9578 non-null uint8  
purpose_home_improvement 9578 non-null uint8  
purpose_major_purchase 9578 non-null uint8  
purpose_small_business 9578 non-null uint8  
dtypes: float64(6), int64(7), uint8(6)  
memory usage: 1.0 MB
```

```
In [15]: from sklearn.model_selection import train_test_split
```

```
In [16]: X = final_data.drop('not.fully.paid', axis=1)  
y = final_data['not.fully.paid']  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3  
0, random_state=101)
```



```
In [17]: from sklearn.tree import DecisionTreeClassifier
```

```
In [18]: dtree = DecisionTreeClassifier()
```

```
In [19]: dtree.fit(X_train,y_train)
```

```
Out[19]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=N
one,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort=False, random_state=N
one,
                                splitter='best')
```

```
In [20]: predictions = dtree.predict(X_test)
```

```
In [21]: from sklearn.metrics import classification_report,confusion_matrix
```

```
In [22]: print(classification_report(y_test,predictions))
```

	precision	recall	f1-score	support
0	0.86	0.82	0.84	2431
1	0.20	0.24	0.22	443
avg / total	0.75	0.73	0.74	2874

```
In [23]: print(confusion_matrix(y_test,predictions))
```

```
[[1996  435]
 [ 336  107]]
```

```
In [24]: from sklearn.ensemble import RandomForestClassifier
```

```
In [25]: rfc = RandomForestClassifier(n_estimators=600)
```

```
In [26]: rfc.fit(X_train,y_train)
```

```
Out[26]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=None, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=600, n_jobs=1,
                                oob_score=False, random_state=None, verbose=0,
                                warm_start=False)
```

```
In [27]: predictions = rfc.predict(X_test)
```

```
In [28]: from sklearn.metrics import classification_report, confusion_matrix
```

```
In [29]: print(classification_report(y_test,predictions))
```

	precision	recall	f1-score	support
0	0.85	1.00	0.92	2431
1	0.59	0.02	0.04	443
avg / total	0.81	0.85	0.78	2874

```
In [30]: print(confusion_matrix(y_test,predictions))
```

```
[[2424   7]
 [ 433  10]]
```

```
In [ ]:
```